

IF GOINGToCrash(): DONT()
RISC-V ARCHITECTURE FOR MOTION PLANNING
ALGORITHMS IN AUTONOMOUS UAVs

A senior design project submitted in partial fulfillment of the requirements for the degree of
Bachelor of Science at Harvard University

Anthony J.W. Kenny
S.B. Candidate in Electrical Engineering

Faculty Advisor: Vijay Janapa Reddi

Harvard University School of Engineering and Applied Sciences
Cambridge, MA

March 1st, 2020
Version: 4.2

Fix Cover Page Formatting

Abstract

This thesis describes a process for accelerating motion planning in autonomous robots through the design of specialised microarchitecture and instruction set architecture. First, it shows the analysis of computational performance of Rapidly-exploring Random Tree (RRT), a sampling-based motion planning algorithm commonly used in autonomous drones. Having identified collision detection as the biggest area of opportunity for improved performance, it describes the process of designing specialized hardware, taking advantage of parallelization, that quickly detects collisions. Finally, it presents how this specialized functional unit can be implemented in a processor, and a RISC-V Instruction Set Architecture (ISA) extension designed to massively reduce the execution time of collision detection.

Rewrite Abstract

Contents

Preface	i
Abstract	i
Table of Contents	ii
List of Acronyms	v
List of Algorithms	vii
List of Figures	viii
List of Tables	ix
1 Introduction	1
1.1 Problem Summary	1
1.1.1 Background	1
1.1.2 Problem Definition	4
1.2 Prior Work	4
1.2.1 Hardware Acceleration	4
1.2.2 RISC-V	6
1.3 Project Overview	8
1.3.1 Proposed Solution	8
1.3.2 Project Specifications	9
1.3.3 Project Structure	9
2 Motion Planning in Software	11
2.1 Motion Planning Background	11
2.1.1 Key Concepts	12
2.1.2 Rapidly-exploring Random Tree	13
2.2 Implementation of RRT	18
2.2.1 Technical Specifications	18
2.2.2 Implementation Design	19
2.2.3 Visualizing Implementation	20
2.3 Analysis of RRT	23
2.3.1 Experimental Methodology	23

2.3.2	Results	24
3	Motion Planning in Hardware	27
3.1	Defining the Collision Detection Unit	27
3.1.1	Edge Collision Function	27
3.1.2	Technical Specifications	27
3.1.3	Performance Specifications	28
3.2	HoneyBee	29
3.2.1	Design	29
3.2.2	Build	30
3.2.3	Measurement and Analysis	31
4	RISC-V Processor	34
4.1	Introduction to the Reduced Instruction Set Computer	34
4.1.1	Instruction Set Architecture	34
4.1.2	RISC Processor Design	34
4.2	RISC-V ISA	34
4.2.1	RV32I	34
4.2.2	Motion Planning Extension	36
4.3	PhilosophyV	36
4.3.1	Baseline Implementation	36
4.3.2	Implementing HoneyBee	36
4.4	Performance Analysis	38
5	Conclusion	39
5.1	Discussion of Results	39
5.2	Evaluation of Success	39
5.3	Future Work	39
	Bibliography	40
	Glossary	42
	Appendices	45
A	Project Repository	46
B	Budget	47
C	RRT Supporting Documentation	48
C.1	Justification of Modelling UAV as Prism	48
C.2	Full Technical Specifications for RRT Implementation	49

C.3	Assessment of Existing RRT Implementations	50
C.4	Implementation of Key RRT Functions	51
C.5	Geometrically Determining Segment-Plane Intersection	54
C.6	Timing Methodology of RRT Analysis	55
C.7	Justification of K:DIM Ratio	57

List of Acronyms

2D 2-Dimensional

3D 3-Dimensional

API Application Programming Interface

ARM Advanced RISC Machine

ASP Application Specific Processor

CISC Complex Instruction Set Computer

CPI Cycles Per Instruction

CPU Central Processing Unit

CSV Comma Seperated File

DOF Degree-of-Freedom

FPGA Field Programmable Gate Array

GPU Graphics Processing Unit

HB-A HoneyBee-A

HDL Hardware Description Language

HLS High Level Synthesis

ISA Instruction Set Architecture

OGM Occupancy Grid Map

PRM Probabalistic Road Map

RISC Reduced Instruction Set Computer

RRT Rapidly-exploring Random Tree

RRT* Rapidly-exploring Random Tree Star

RTOS Real-Time Operating Systems

RV32I RISC-V 32-Bit Integer

SoC System on Chip

UAV Unmanned Aerial Vehicle

List of Algorithms

2.1	Rapidly-Exploring Random Tree in Free configuration Space	14
2.2	Rapidly-Exploring Random Tree with Collision Detection	17
C.1	getRandomConfig() as implemented for RRT	51
C.2	findNearestConfig() as implemented for RRT	51
C.3	stepFromNearest() as implemented for RRT	52
C.4	configCollision() as implemented for RRT	52
C.5	configCollision() as implemented for RRT for 3D	53

List of Figures

1.1	Simple Visualization of Computer Implementation Hierarchy	5
1.2	Typical Process of Adding Non-Standard Extension to RISC-V ISA	7
1.3	System Diagram of Overall Project	8
2.1	Example of 2 Robot Configurations in 3D Space for Motion Planning Purposes	12
2.2	Occupancy Grid Maps for (16×16) Workspace of Different Resolutions . .	13
2.3	Scope of the RRT Algorithm	14
2.4	Step by step demonstration of RRT Algorithm for 2D robot in 2D space . .	16
2.5	Demonstration of the 5 Key Functions that Constitute RRT	20
2.6	Visualization of Workspace in 2D and 3D	21
2.7	Visualization of Obstacles in 2D and 3D	22
2.8	Complete Visualization of RRT in 2D and 3D	22
2.9	Increasing Total Execution Time of RRT with Map Size	24
2.10	Profile of Computational Load of RRT in 2D	25
2.11	Profile of Computational Load of RRT in 3D	26
3.1	Pipelining to Improve Latency	29
3.2	Interface Summary of HoneyBee-A Synthesis in Vivado HLS	31
3.3	RRT Simulated Execution Time with HB-A (microseconds)	33
4.1	5-Stage Reduced Instruction Set Computer (RISC) Datapath	34
4.2	Philosophy V Processor	37
C.1	Modelling a UAV as a Rectangular Prism	48
C.2	Using Parallel Planes to determine Edge Collisions with Grids	54

List of Tables

1.1	List of System Components and their Descriptions	9
2.1	Abbreviated Technical Specifications for RRT Implementation	18
2.2	RRT Implementation Parameters	19
2.3	Optimal RRT Parameters for each Map Size	24
3.1	Technical Specifications for Edge Collision Detection Unit	28
3.2	Simulated performance of HB-A in microseconds	32
3.3	RRT Simulated Execution Times with HB-A (seconds)	32
4.1	Register State for RV32I Base Instruction Set	35
4.2	RV32I Base Instruction Formats	36
C.1	General Technical Specifications for RRT Implementation	49
C.2	Required Parameters for RRT Implementation	50
C.3	Evaluation of Existing Open-Source Implementations of RRT	50
C.4	Comparison of Timing Methods	56

Chapter 1

Introduction

Define goal: Somewhere here I need to define very clearly the following

1. Overall Goal:~ Deliver an example of how RISC-V can be leveraged to design a custom processor for motion planning
2. 4 Objectives:
 - Profile a standard motion planning algorithm to determine the bottleneck function
 - Design a functional hardware unit that eliminates that bottleneck
 - Define a non standard extension
 - Implement Processor and verify all above

1.1 Problem Summary

1.1.1 Background

The Unmanned Aerial Vehicle (UAV) has been utilised in military applications extensively throughout the late 20th and early 21st century. However, over the last decade, their use in non-military applications, such as commercial, scientific, agricultural, and recreational, has increased such that the number of civilian drones vastly outnumber military UAVs. Particularly in the commercial sector, such rapid growth in the number and range of applications means that autonomy is key for the profitable adoption of UAVs. Such autonomy relies on efficient computation of motion planning algorithms. However, the implementation of these algorithms can be quite computationally expensive, and thus slow and/or detrimentally power consuming. As such, this thesis aims to design specialized hardware to more efficiently compute motion plans for autonomous drones.

cite

Autonomous Robotics

For well over 2000 years, the concept of robotics, albeit not always with such a term, has fascinated humans. As early as the first century A.D., the Greek mathematician and engineer, Heron of Alexandria, described more than 100 different machines and automata in *Pneumatica* and *Automata* [1]. In 1898, Nikola Tesla demonstrated the first radio-controlled vessel. Since then, the world has seen widespread application of robotics in manufacturing, mining, transport, exploration, and weaponry. For the last few decades, robots have operated in controlled, largely unchanging environments (e.g. an assembly line) where their environment and movements are largely known *a priori*.

However, in recent years a new generation of autonomous robots has been developed for a wide range complex applications. These new robots are required to adapt to the changing environments in which they operate; most often, this means planning their own paths through space. As such, they must perform motion planning in real-time.

Motion Planning

While most creatures in the animal kingdom find it relatively easy to navigate their surroundings, autonomous robots must be taught explicitly how to do so by their programmers. Motion Planning refers to the problem of algorithmically determining a collision-free path between two points in an obstacle-ridden space. Chapter 2 provides a detailed explanation of motion planning and of Rapidly-exploring Random Tree (RRT), a commonly used motion planning algorithm.

On the algorithmic and software level, motion planning has been extensively studied and optimized. Even so, current software implementations running on regular Central Processing Units(CPUs) are too slow to execute in real-time for robots to operate in rapidly changing, high complexity environments. More powerful, highly parallelized Graphics Processing Units(GPUs) can be used in tethered robot applications (e.g. robotic arms autonomously executing pick-and-place functions). However, such GPUs consume far too much power to be used in autonomous drones, which are untethered and must sustain flight for useful periods of time. (A typical CPU uses between 65-85 watts, while some GPUs can use up to 270 watts).

cite

Application Specific Processors

Given the lacking performance in computing motion plans of a CPU, and the untenable power consumption of a GPU, autonomous drone developers are left with the option of developing an Application Specific Processor (ASP), optimized for motion planning.

However, designing a functional, high performance processor from scratch is no small task. It requires expertise in a variety of disciplines (compilers, digital logic, operating systems, etc), and an extraordinary amount of time and effort to develop and verify before it can be used. In short, it's an expensive process, which is why the market for computer

processors is dominated by companies like Intel, AMD, and ARM. The sharing of processor designs is also not possible, as commercial designs are proprietary and competing designs are not encouraged.

Finally, even if one were to design an ASP from scratch, or build off an existing commercial design (which means paying royalties), the Instruction Set Architecture (ISA) that the processor implements are not designed for extendability, meaning that even a highly specialized processor is limited to a small number of instructions.

Discuss Moore's law and denard scaling

RISC-V

RISC-V (pronounced “risk-five”) is an ISA developed by the University of California, Berkeley. It is established on the principles of a RISC, a class of instruction sets that allow a processor to have fewer Cycles Per Instruction (CPI) than a Complex Instruction Set Computer (CISC) (x86, the ISA on which macOS and linux operating systems run, is an example of a CISC instruction set).

What makes RISC-V unique is its open-source nature. RISC-V was started with the philosophy of creating a practical, open-source ISA that was usable in any hardware or software without royalites. The first report describing the RISC-V Instruction Set was published in 2011 by Andrew Waterman, Yunsup Lee, David A. Patterson, and Krste Asanović [2].

1.1.2 Problem Definition

Problem Statement

Motion planning algorithms implemented in software that runs on general purpose CPUs cannot execute quickly enough for fully autonomous UAVs to operate in high-complexity environments. The state-of-the-art strategy of using power-hungry GPUs to accelerate the execution of these algorithms requires too much power to be cost-effective or feasible for UAVs to sustain flight for useful periods of time.

Improve Problem Statement: Existing research into accelerating robotic motion planning is <reason for RISC-V, inaccessible?> and mainly focussed on tethered arm moving robots.

End User

This thesis aims to provide developers of autonomous drones with specialized hardware for motion planning. Such developers have a need for computing hardware that executes motion planning algorithms faster and more power efficiently than existing methods. This thesis will provide a processor design that is synthesizable on an Field Programmable Gate Array (FPGA), giving developers a processor for which a Real-Time Operating Systems (RTOS), or bare metal code, can be deployed.

Revise End User

1.2 Prior Work

1.2.1 Hardware Acceleration

Hardware acceleration refers to the strategy of using computer hardware specifically designed to execute a function more efficiently than can be achieved by software running on a general purpose CPU. Specialized hardware designed to perform specific functions can yield significantly higher performance than software running on general purpose processors, and lower power consumption than GPUs.

Computer Implementation Hierarchy

To briefly frame the space in which this thesis operates, consider the typical computer implementation hierarchy, demonstrated in Figure 1.1. **User level applications**, such as Google Chrome, Microsoft Word, and Apple's iTunes, sit at the top of the abstraction hierarchy. These applications are implemented in **High-Mid Level Languages**, such as C/C++, Python, Java, etc. These programming languages have their own hierarchy, but for the purpose of this thesis, it is sufficient to understand that these programming languages are then compiled into **Assembly Language**. Assembly language closely follows the execution of instructions on the **processor**, and is defined by an **ISA**. An ISA can be

thought of as the contract between software programmers and processor engineers, agreeing what instructions the processor is able to implement. This assembly code is finally loaded into the processor's instruction memory and executed.

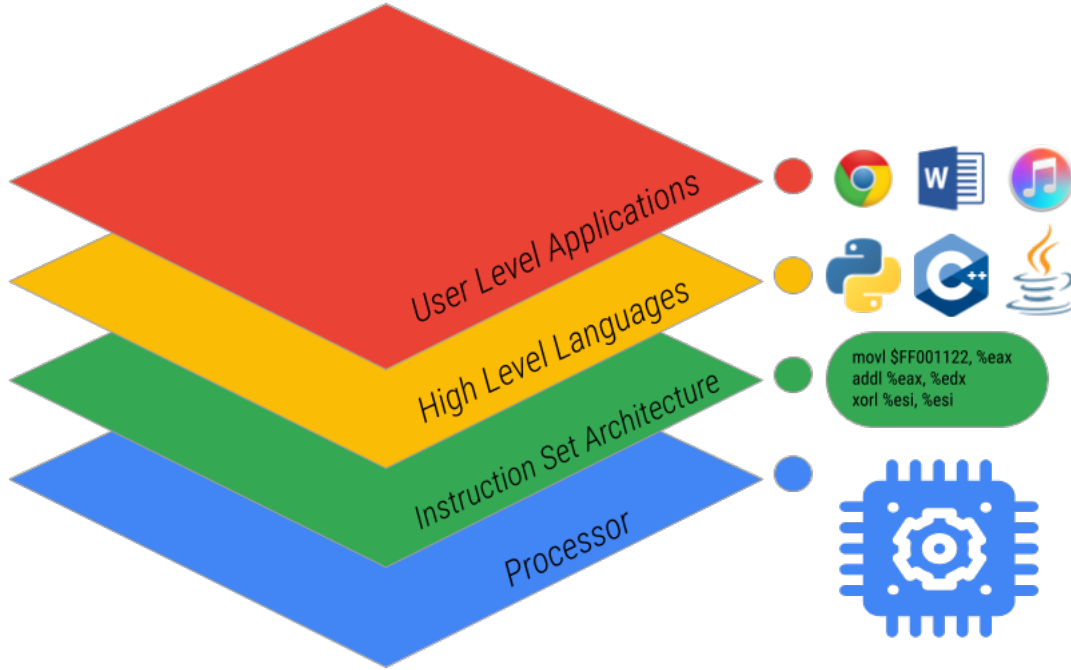


Figure 1.1: Simple Visualization of Computer Implementation Hierarchy

As will be outlined in Section 1.3, this thesis operates extensively on the lower two levels of this hierarchy, extending an existing ISA and building hardware at the processor level that supports these extensions.

Acceleration of Motion Planning

Accelerating motion planning with hardware is a fairly well studied problem.

A Motion Planning Processor on Reconfigurable Hardware [3] studied the performance benefits of using FPGA-based motion planning hardware as either a motion planning processor, co-processor, or collision detection chip. It targeted the feasibility checks of motion planning (largely collision detection) and found their solution could build a roadmap using the Probabilistic Road Map (PRM) algorithm up to 25 times faster than a Pentium-4 3Ghz CPU could.

In *A Programmable Architecture for Robot Motion Planning Acceleration* [4], Murray et

al. built on the work of the aforementioned paper, to accelerate several aspects of motion planning in an efficient manner.

FPGA based Combinatorial Architecture for Parallelizing RRT [5] studies the possibility of building architecture to allow multiple RRTs to work simultaneously to uniformly explore a map. Taking advantage of hardware parallelism allows systems such as this to compute more information per clock cycle.

Finally, in the paper *Robot Motion Planning on a Chip* [6], Murray et al. describe a method for constructing robot-specific hardware for motion planning, based on the method of constructing collision detection circuits for PRM that are completely parallelised, such that edge collision computation performance is independent of the number of edges in the graph. With this method, they could compute motion plans for a 6-degree-of-freedom robot more than 3 orders of magnitude faster than previous methods.

1.2.2 RISC-V

Extending RISC-V

RISC-V is designed cleverly in a modular way, with a set of base instruction sets and a set of standard extensions. As a result, processors can be designed to only implement the instruction groups it requires, saving time, space and power on instructions that won't be used. In addition, another goal of RISC-V is to provide a basis for more specialized instruction-set extensions or more customized accelerators. This is described in the most recent *RISC-V Instruction Set Manual* [7]. This is a powerful feature, as it does not break any software compatability, but allows for designers to easily follow the steps outlined in Figure 1.2. From a hardware acceleration point of view, this is particularly useful as it allows the designer to directly invoke whatever functional unit or accelerator they implement from assembly code.

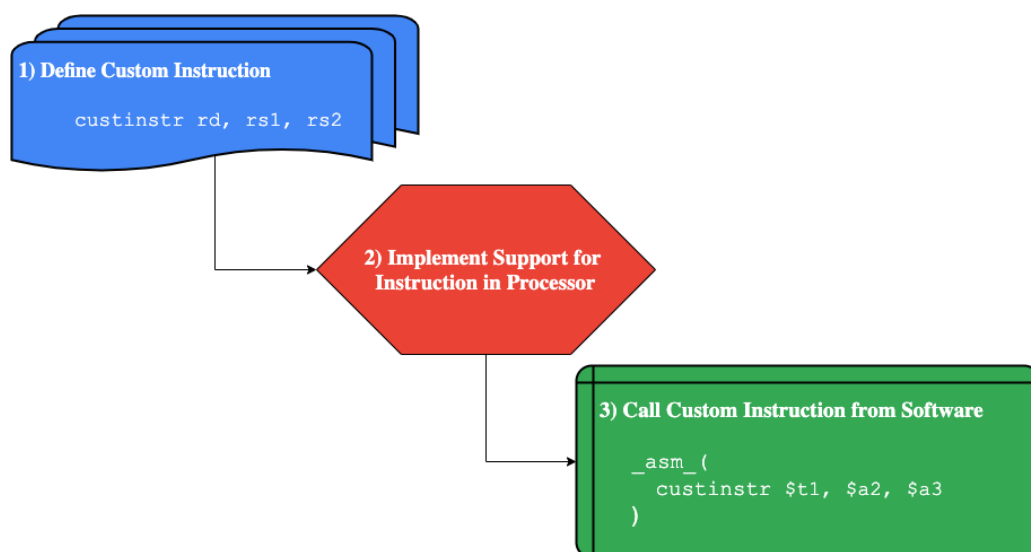


Figure 1.2: Typical Process of Adding Non-Standard Extension to RISC-V ISA

Accelerating RISC-V Processors

Having only been released in 2011, RISC-V is still a relatively unexplored opportunity for non-education applications. However, it shows promise in the commercial space, with Alibaba recently developing the Xuantie, a 16-core, 2.5GHz processor, currently the fastest RISC-V processor. Recently there has been promising research into accelerating computationally complex applications, particularly in edge-computing, with RISC-V architecture. *Towards Deep Learning using TensorFlow Lite on RISC-V*, a paper co-written by the faculty advisor of this thesis, V.J. Reddi, presented the software infrastructure for optimizing the execution of neural network calculations by extending the RISC-V ISA and adding processor support for such extensions. A small number of instruction extensions achieved coverage over a wide variety of speech and vision application deep neural networks. Reddi et al. were able to achieve an 8 times speedup over a baseline implementation when using the extended instruction set. *GAP-8: A RISC-V SoC for AI at the Edge of the IoT* proposed a programmable RISC-V computing engine with 8-core and convolutional neural network accelerator for power efficient, battery operated, IoT edge-device computing with order-of-magnitude performance improvements with greater energy efficiency.

1.3 Project Overview

1.3.1 Proposed Solution

This thesis proposes a non-standard RISC-V Instruction Set Extension, supported by a functional unit embedded in a FPGA synthesizable processor design that more rapidly computes motion planning for autonomous UAVs. It will use the RRT algorithm as a benchmark for performance analysis. Profiling of RRT described in chapter 2 found that edge collision detection was the most performance limiting function of RRT. As such, this thesis aims to design a RISC-V extension and specific circuitry that support the faster execution of edge collision detection.

System Overview

Figure 1.3 shows a high level overview of the system this thesis proposes.




Figure 1.3: System Diagram of Overall Project

The **Extended RISC-V ISA** is made up of the RISC-V 32-Bit Integer (RV32I) Base Instruction Set and a non-standard extension that this thesis will define. The **PhilosophyV Processor** is a RISC-V chip built in Hardware Description Language (HDL) for this thesis. It implements both the RV32I instruction set and the non-standard extension. The PhilosophyV Core includes, along with a baseline 5-stage processor implementation, the **HoneyBee** collision detection unit. A **C Implementation of RRT** is loaded into the instruction memory of the PhilosophyV processor. This processor, synthesized on an **FPGA**, is used as the main processor, co-processor, or accelerator on an **Autonomous UAV**. Table 1.1 outlines the components of this system and their descriptions.

Component	Source	Description
RISC-V Instruction Set		
RV32I	Berkeley	40 Instructions defined such that RV32I is sufficient to form a compiler target and suport modern operating systems [7].
Extension	<i>New</i>	This is the custom extension defined by this thesis targeting motion planning instructions. It is outlined in Chapter 4.
C-Implementation of RRT		
RRT	<i>New</i>	Due to lack of available implementations of RRT suitable for the purposes of this thesis, RRT was implemented from the ground up in C. This is detailed in Chapter 2
FPGA Synthesized Chip		
Zynq-7000	Xilinx	The Zynq-7000 family of System on Chip (SoC)s are a low cost FPGA and Advanced RISC Machine (ARM) combined unit.
PhilosophyV	<i>New</i>	The processor built for this thesis to demonstrate how the RISC-V extension and hardware unit work together. This is detailed in Chapter 4
HoneyBee	<i>New</i>	The functional unit designed specifically for faster execution of edge collision detection computations. Outlined in Chapter 3

Table 1.1: List of System Components and their Descriptions

1.3.2 Project Specifications

Project Specifications: These need significant revision from the last checkpoint

1.3.3 Project Structure

This report is structured to follow the timeline of this project, and is outlined below:

1. A benchmark motion planning algorithm, RRT, was selected and implemented in software. Once implemented, a variety of performance analysis methods were used to profile the computational hotspots of the algorithm. It was found that edge collision detection was the critical function limiting execution time. This process is detailed in Chapter 2.

2. With edge collision detection having been identified as the critical function, the process of designing specialised hardware to execute this function began. The technical specifications, performance specifications, designs, build phases, measurement and analysis of this hardware unit is presented in Chapter 3.
3. With the aforementioned functional unit's performance verified in simulation, the next step was to implement this in a processor. First, a baseline processor was designed and built for this project to implement a base RISC-V instruction set. The performance of RRT is again profiled on this baseline processor (as up until this point, it was profiled on x86 architecture). A non-standard extension to the RISC-V ISA was then defined and support for this was implemented in the processor. Comparative performance analysis was then conducted. This process is described in detail in Chapter 4.
4. Chapter 5 is a discussion of results and future work.

Summary of Results: Do I need a summary of results section in the introduction?

Chapter 2

Motion Planning in Software

The first objective of this thesis is to identify a typical motion planning algorithm, profile its execution, and determine computational bottlenecks.

This chapter introduces the concept of motion planning and details the process of implementing and analysing Rapidly-exploring Random Tree (RRT), a commonly used algorithm, to identify its computational bottlenecks.

Update
once
I have
properly
defined
goals and
objectives

2.1 Motion Planning Background

A funny paradox in computer science is the fact that it is relatively easy to teach a computer to perform tasks that humans find very complicated, but extremely difficult to program one to execute functions that humans master during infancy. Consider, it was as early as 1949 that Claude Shannon presented his paper *Programming a Computer for Playing Chess*[8], and by 1997 the *Deep Blue* computer defeated Garry Kasparov, reigning world champion, in a six game chess match.[9] Compare that with some of the most advanced autonomous humanoid robots to date displaying dexterity only comparable with that of a toddler. The task of finding a collision free path, performed constantly without thought by a human, is an example of this paradigm. For a robot to compute a collision free path, it relies on a set of Motion Planning Algorithms.

Motion Planning Algorithms refer to the set of algorithms that find possible sequences of valid configurations for a robot in a space. In plain English, they are algorithms that determine the movements a robot can make in a map, with the intent of eventually finding a path from one point to another.

2.1.1 Key Concepts

Workspace

The workspace, more loosely known as the **map**, is the space which the robot and obstacles occupy. Obviously, **obstacles** refer to anything with which the robot cannot intersect.

Configuration

A configuration describes the position, orientation, and pose of the robot. The complexity of a robot's configuration is therefore dependant on the dimension of the workspace, the complexity of the robot itself, and in what level of detail the robot must be represented. For example:

- Most simply, a robot can be represented as a point by the Cartesian coordinates (x, y) 2-Dimensional (2D) space and (x, y, z) in 3-Dimensional (3D) space.
- More realistically, a robot such as a drone may be represented in 3D by an origin point (x, y, z) and 3 Euler angles (α, β, γ) describing its orientation.
- In a more complex form, a fixed base robot with N Degree-of-Freedom (DOF) would require an N -dimensional configuration.




Figure 2.1: Example of 2 Robot Configurations in 3D Space for Motion Planning Purposes

Occupancy Grid Map

An Occupancy Grid Map (OGM) is a method of representing the obstacles present in a workspace. Obstacles are often irregularly shaped and computing collisions with such obstacles is near impossible. Therefore, the workspace is discretized into grids and grids containing any part of the obstacle are marked as occupied, even if only a small part of the grid is occupied. An Occupancy Grid Map (OGM) will more accurately represent a workspace with a higher resolution, shown in Figure 2.2.

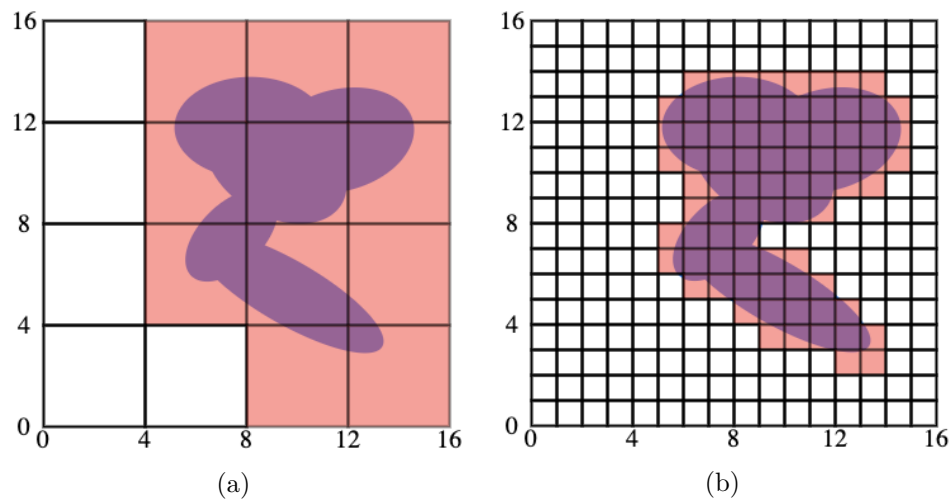


Figure 2.2: Comparison of Occupancy Grid Maps for 16×16 Workspace with Different Resolutions. Figure 2.2a shows how an OGM with low resolution, while simpler to construct and analyse, will over-represent the obstacle density of a workspace. Figure 2.2b shows how a higher resolution will more accurately reflect the obstacles of a workspace.

2.1.2 Rapidly-exploring Random Tree

Rapidly-exploring Random Tree (RRT) is an algorithm designed to efficiently build a tree of collision-free paths in a high-complexity environment. The algorithm randomly samples points, draws an edge from the nearest currently existing node in the tree, to grow the tree in the space. It is inherently biased to grow towards large unsearched areas of the workspace. RRT was developed by S. LaVelle[10] and J. Kuffner[11]. It is used in autonomous robotic motion planning problems such as autonomous drones.

Scope

RRT takes an initial configuration, a goal point, and an Occupancy Grid Map (OGM) as its input. This OGM may be built and updated using a priori knowledge, sensor data from

the robot, and other inputs. RRT will output a tree of collision free paths toward the goal, as demonstrated in Figure 2.3. **It does not calculate the fastest path from that tree;** that can be accomplished using algorithms such as Dijkstra's algorithm.




Figure 2.3: Scope of the RRT Algorithm: Takes an OGM as input and outputs a tree of collision free paths through that OGM. The tree is shown in blue on the right.

Algorithm

Put simply, RRT finds a path from start to finish by randomly exploring a workspace. Put more technically, it builds a tree of possible configurations (also known as a graph), connected by edges, for a robot of some physical description. It does so by selecting random configurations and adding them to the graph. From this graph, a path from the initial configuration to some goal configuration can be found, given a high enough number of iterations. As such, RRT can be considered probabilistically complete. The pseudo-code for RRT can be seen in Algorithm 2.1

Algorithm 2.1: Rapidly-Exploring Random Tree in Free configuration Space

Inputs: Initial configuration q_{init} ,
Number of nodes in graph K ,
Incremental Distance Δq

Output: RRT Graph G with K configurations $[q]$ & edges $[e]$

```

G.init();
for  $k = 1$  to  $K$  do
     $q_{rand} \leftarrow \text{randomConfiguration}()$ ;
     $q_{near} \leftarrow \text{nearestVertex}(q_{rand}, G)$ ;
     $q_{new} \leftarrow \text{newVertex}(q_{near}, q_{rand}, \Delta q)$ ;
     $G.\text{addVertex}(q_{new})$ ;
     $G.\text{addEdge}(q_{near}, q_{new})$ ;
end

```

Algorithm 2.1 can be visually represented in Figure 2.4. Consider a 2D robot operating in a 2D workspace. A Graph G is initialized containing an initial configuration, q_{init} , with constraints on the number of nodes that the graph can hold, K , and the maximum distance between two nodes, Δq . This is shown in Sub-figure 2.4a. A random configuration for the robot, q_{rand} is generated (2.4b). The nearest existing configuration in G , q_{near} , is found. (In the first iteration, $q_{near} = q_{init}$, shown in Sub-figure 2.4c). The distance between q_{near} and q_{rand} is calculated. If this distance is less than Δq , $q_{new} = q_{rand}$. If not, q_{new} is selected, typically by moving by Δq from q_{near} towards q_{rand} (2.4c). q_{new} is then added to G . This is repeated for K configurations.

This is all really ugly and should be explained better

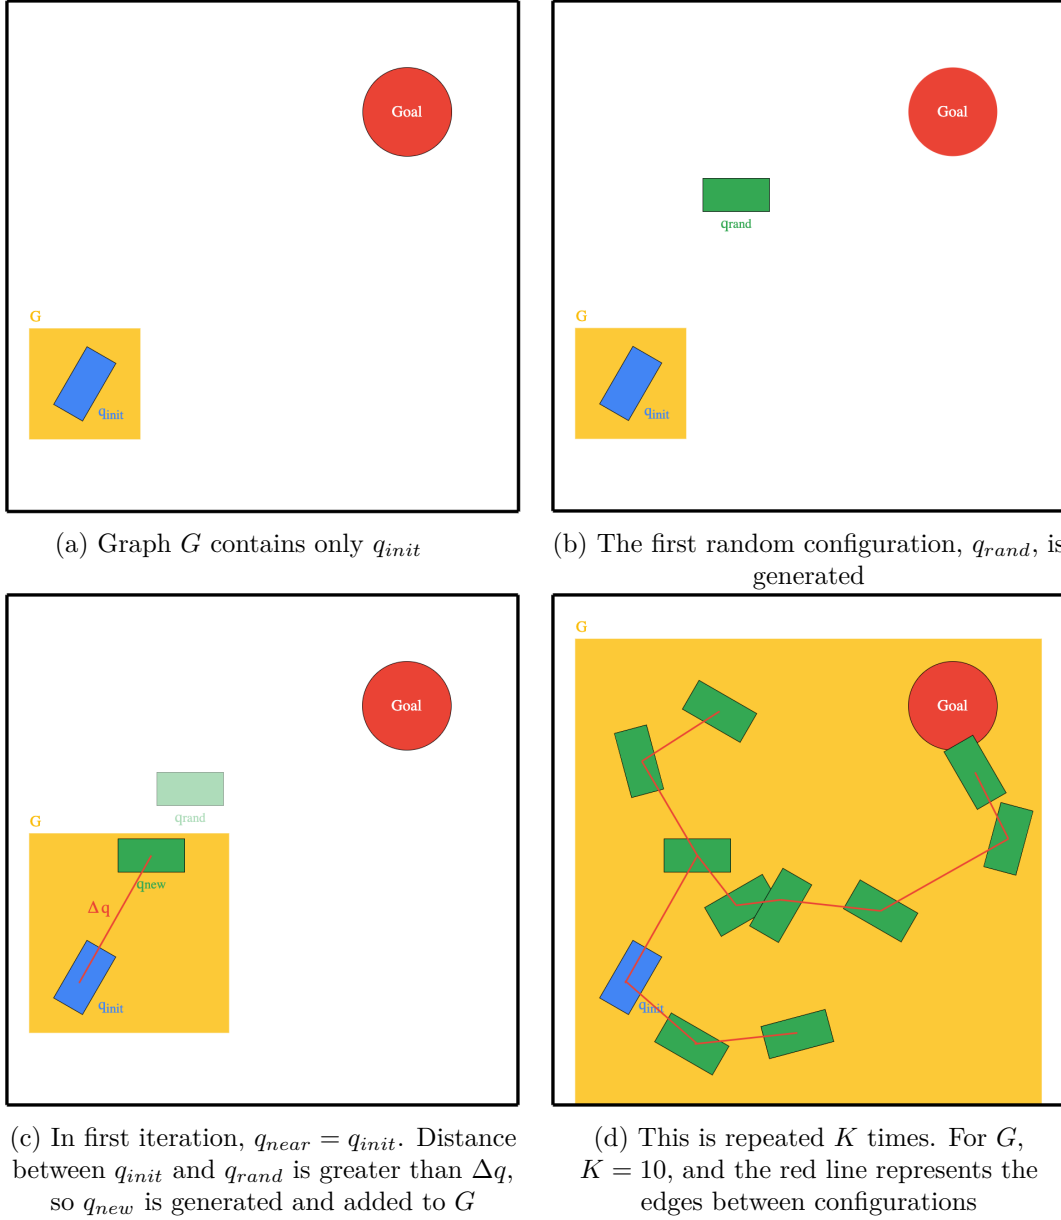


Figure 2.4: Step by step demonstration of RRT Algorithm for 2D robot in 2D space

Algorithm 2.1 shows how RRT builds a graph of possible configurations connected by edges in a completely free configuration space. However, in real-world applications,

Redo this diagram

a robot's workspace space often contains obstacles. As such, collision detection must be included in the algorithm. The two types of collisions the algorithm must check for are *configuration collisions* (those where the robot would collide with an obstacle in a given configuration) and *edge collisions* (where the robot would collide when moving between two collision free configurations).

The RRT with configuration and edge collision detection can be seen in Algorithm 2.2. The method of implementing RRT with collision detection to model a drone in 3D space is detailed in Section 2.2.

Algorithm 2.2: Rapidly-Exploring Random Tree with Collision Detection

Inputs: Initial configuration q_{init} ,
Number of nodes in graph K ,
Incremental Distance Δq ,
Space S containing obstacles

Output: RRT Graph G with K configurations $[q]$ & edges $[e]$

```

G.init();
for  $k = 1$  to  $K$  do
    while !pointCollision( $q_{new}$ ) do
         $q_{rand} \leftarrow \text{randomConfiguration}()$ ;
         $q_{near} \leftarrow \text{nearestVertex}(q_{rand}, G)$ ;
         $q_{new} \leftarrow \text{newVertex}(q_{near}, q_{rand}, \Delta q)$ ;
    end
     $e_{new} \leftarrow \text{newEdge}(q_{near}, q_{new})$ 
    if !edgeCollision( $e_{new}$ ) then
        G.addVertex( $q_{new}$ );
        G.addEdge( $q_{near}, q_{new}$ );
    else
         $k = k - 1$ ;
    end
end

```

2.2 Implementation of RRT

2.2.1 Technical Specifications

With RRT selected as the benchmark algorithm against which to test specialised hardware, this project required an implementation of the algorithm that satisfied the following criteria shown in Table 2.1. Appendix C.2 is a more thorough description of the technical specifications for the implementation of RRT.

Requirement	Brief Description and Justification
Implemented in C/C++	Implementations in C allow for more accurate analysis of computational bottlenecks, unlike higher-level languages like Python.
3D Workspace	The computational requirements of RRT in 3D differ somewhat to that in 2D. Since autonomous UAVs operate in 3D space, it was necessary to have a 3D implementation to analyse.
UAV modelled as a 3D rectangular prism	In theory, it is possible to model a UAV much more precisely than a rectangular prism. However, in reality, modelling a UAV as a 3D rectangular prism, defined by coordinates $\{x, y, z\}$ and Euler angles $\{\alpha, \beta, \gamma\}$, is more than sufficient (and more computationally efficient). See Appendix C.1 for justification.
Mathematically Complete Collision Detection	When RRT is implemented for educational purposes, the edge collision calculations are often simplified to a sampling model which is probabilistically complete but not mathematically complete. In other words, it will catch most collisions by sampling a number of points along each edge, but there is always a possibility of an undetected collision. In real world applications, collisions must be calculated by method of geometric intersection to ensure all collisions are detected.
Highly Parameterizable	Accurate analysis of the algorithm required the ability to vary the following parameters: <ul style="list-style-type: none"> • ϵ (Maximum distance between two configurations) • K (Maximum number of configurations) • DIM (The upper bound of each dimension for a $DIM \times DIM \times DIM$ workspace) • Goal Bias (How biased RRT is to move towards goal point)

Table 2.1: Abbreviated Technical Specifications for RRT Implementation

The original intention was to find an existing implementation of RRT that could fulfill these requirements. However, no open-source implementations were suitable. Appendix C.3 shows an evaluation of existing implementations.

As a result, it was necessary to build a C implementation of RRT from the ground up to the aforementioned specifications.

2.2.2 Implementation Design

The design and implementation of RRT, while necessary, was significant and time consuming. Since this was not the main object of this thesis, only a brief description of key design choices has been included here. Appendix C contains a more detailed account.

System
Diagram
Here

Parameterization

Table 2.2 shows the parameters that were included in the implementation and compiled by way of a C header file.

Parameter	Data Type	Description
ϵ	Integer	Maximum distance between two configurations
K	Integer	Maximum number of configurations in the graph
DIM	Integer	Upper bound of each axis of workspace
Goal Bias	Float	Percentage likelihood of stepping towards goal node
OGM	File Pointer	CSV of booleans to represent grids

Table 2.2: RRT Implementation Parameters

Dimensionality

RRT was implemented in both 2D and 3D. Not only did a 2D implementation provide a good development checkpoint, it was also interesting to see the difference in computational load between 2D and 3D, discussed in Section 2.3.

Modelling a UAV

The UAV was modelled as a 3D rectangular prism, with its configuration represented by the Cartesian coordinates (x, y, z) and the Euler angles (α, β, γ) .

Key Functions

Algorithm 2.2 shows that there are 5 key functions that constitute RRT. Figure 2.5 demonstrates each of these functions: `getRandomConfig()`, `findNearestConfig()`, `stepFromNearest()`, `(configCollisions)`, and `edgeCollisions()`. Appendix C.4 shows in detail how each of these functions was implemented.

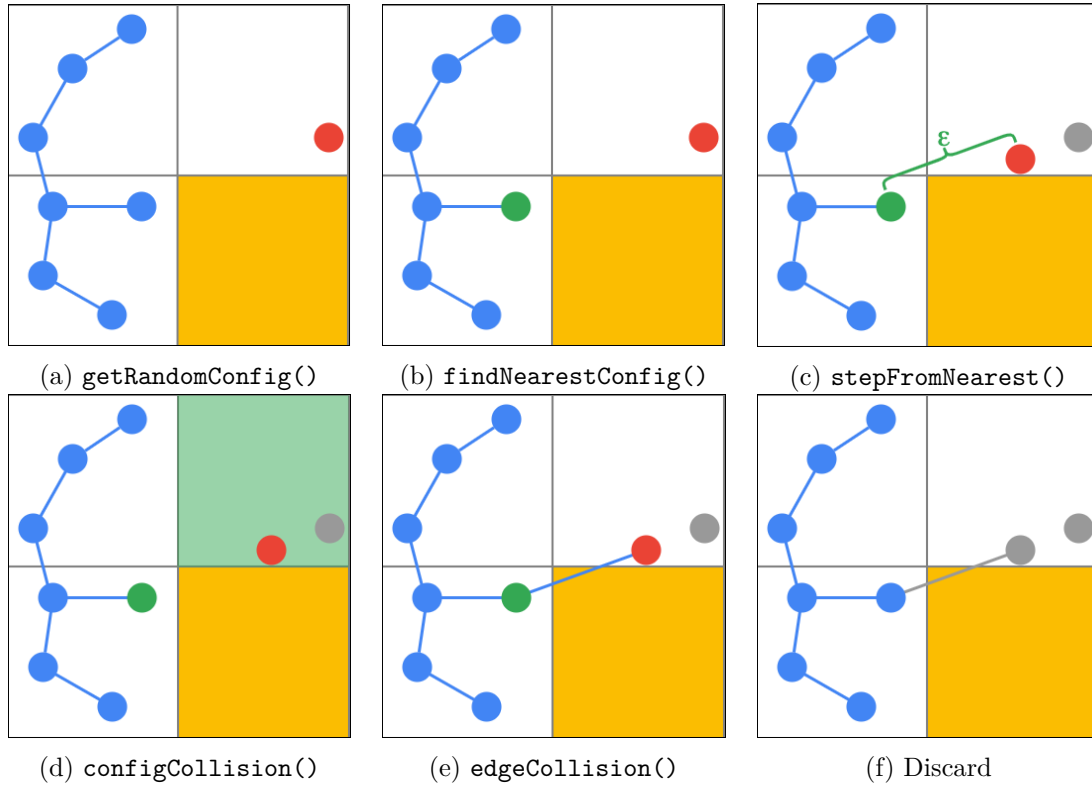


Figure 2.5: **Demonstration of the 5 Key Functions that Constitute RRT**, where configuration is a node in a 2D workspace. A new node is generated in (a) with `getRandomConfig()`, and the closest existing node is found with `findNearestConfig()` in (b). In this case, the new node is further than ϵ from the nearest node, and so a new node is generated with `stepFromNearest()` in (c). `configCollision` determines that the new node is not in an occupied grid (d) and draws an edge between the two nodes. `edgeCollision()` determines that, in this case, there is a collision (e) and the new node is discarded (f).

2.2.3 Visualizing Implementation

With the back end functionality of RRT designed and implemented, it was necessary to develop a way to visualize it. Many existing implementations had the visualization interface run synchronously alongside RRT. This would distort any performance analysis results, and so in this implementation it was left until after RRT had finished executing and then plotted using Python.

Plotting Configurations and the Workspace

Plotting the workspace using the “matplotlib” library was relatively simple in both 2D and 3D, shown in Figure 2.6. It was decided that the UAV’s’ configuration would be visualized only as its origin point, rather than plotting a 3D rectangular prism at each configuration, in order to maintain simplicity. Nevertheless, the UAV was still modelled as a 3D prism in the backend.

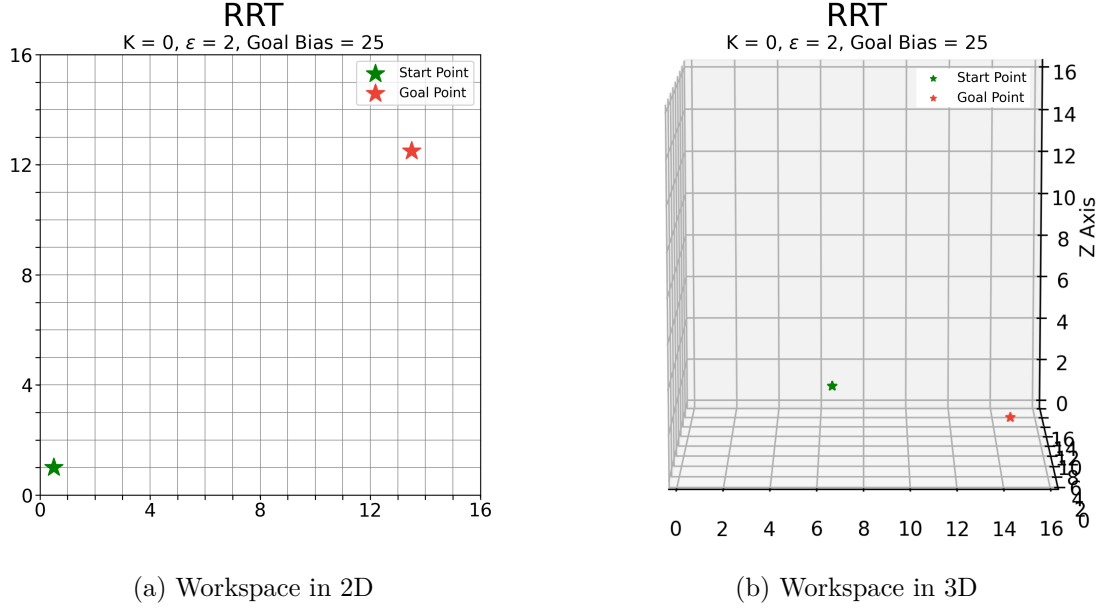



Figure 2.6: **Visualization of Workspace in 2D and 3D**, with configuration represented by only a point, and Start and Goal nodes shown

Plotting Obstacles


Obstacles were plotted in accordance to the input OGM, shown in Figure 2.7

Plotting RRT Graph

To keep the plot simple, it was decided to not show the origin point of each configuration in the graph produced by RRT. Instead, only the edges of the graph were plotted, seen in Figure 2.8




(a) Obstacles in 2D

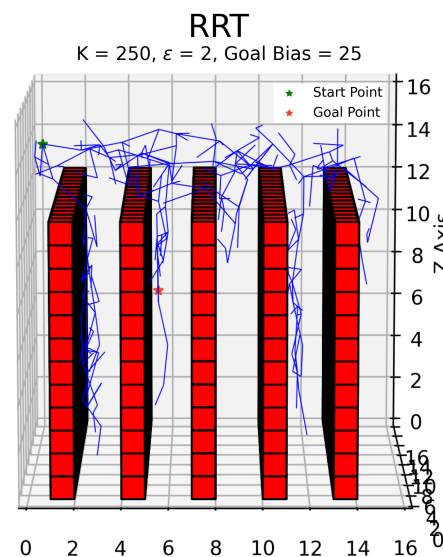


(b) Obstacles in 3D

Figure 2.7: **Visualization of Obstacles in 2D and 3D**, Obstacles shown in yellow and red for 2D and 3D respectively.



(a) RRT in 2D



(b) RRT in 3D

Figure 2.8: **Visualization of Obstacles in 2D and 3D**, with Graph shown in blue.

2.3 Analysis of RRT

Having implemented a functioning version of RRT that adhered to the specifications set out in Table 2.1, analysis of its computational profile could begin. The purpose of this analysis is to identify the biggest bottleneck of RRT and therefore the best opportunity for hardware acceleration.

2.3.1 Experimental Methodology

Experiments were set up to determine which of the 5 key functions of RRT takes up the biggest share of computational load. The only fair way of determining computational load of each function is to measure the percentage of CPU time each function takes in the **fastest possible execution** of RRT for a given map size. This is explained in more detail in Section 2.3.1.

Measuring Performance

The “performance” metric of interest is the percentage of total time the CPU spends executing each of the 5 aforementioned key functions. CPU analysis of a program can often be more complicated than merely timing how long each function takes to execute. Software can be written with inbuilt multithreading and other optimizations that require special CPU analysis software, such as Intel’s VTune Profiler[12]. This software is designed to find computational bottlenecks in large, complex programs. However, it takes significantly longer to run (which was unsuitable for running hundreds of thousands of tests) and is less customizable than adding performance timers directly to the program’s code. It was also hypothesized that, since this project’s implementation of RRT did not use multithreading or any other timing distorting optimizations, custom performance tracking should yield the same results as VTune Profiler. As such, custom performance tracking was added to the RRT implementation. This custom performance tracking method was verified by conducting a χ^2 test against data from VTune Profiler, and was found to be accurate. Appendix C.6 gives more detail on timing methodology.

Optimal Parameters

Extensive testing was undertaken to determine the optimal parameters for a given map size. The goal was to find the set of parameter values for which RRT would reach its goal with $\geq 98\%$ probability, for a wide variety of OGMs, in the shortest possible time.

For each map size 4, 8, 16, 32, 64, the parameters that were varied were ϵ , K , and Goal Bias. The success rate and average execution time was measured by, for each set of parameter values, running RRT 100 times. Thus, with 5 different map sizes, if 4 values were tested for each parameter, and 4 different OGMs were tested, the total number of tests = $4^4 \times 5 = 1280$ (with each test running RRT 100 times!)

As such, only the optimal parameter values for each map size in included in Table 2.3.

DIM	K	ϵ	Goal Bias (%)	Success Rate (%)
2D				
4×4	75	1	10	100
8×8	100	2	25	98
16×16	125	4	25	99
32×32	250	8	10	100
64×64	500	16	25	100
3D				
$4 \times 4 \times 4$	75	1	10	99
$8 \times 8 \times 8$	100	2	25	100
$16 \times 16 \times 16$	100	4	25	100
$32 \times 32 \times 32$	250	8	10	99
$64 \times 64 \times 64$	500	16	25	100

Table 2.3: **Optimal RRT Parameters for each Map Size**, shows the optimal set of parameters after extensive testing, alongside their respective success rates over 100 executions of RRT for different OGMs.

2.3.2 Results

As expected, the total execution time of RRT for optimal parameters increased as the size of the map increased, shown in Figure 2.9.




Figure 2.9: **Increasing Total Execution Time of RRT with Map Size**

2D Computational Load Profile

Figure 2.11 shows that the two biggest computational loads are `findNearestConfig()` and `edgeCollisions()`, with the latter increasing as the size of the map increases. The fact that the load of `edgeCollisions()` takes the majority of execution in bigger map sizes means that, at least in 2D, it can be considered the bottleneck function.




Figure 2.10: Profile of Computational Load of RRT in 2D

3D Computational Load Profile

The computational load of `edgeCollisions()` was even greater in 3D, starting at 40% for $4 \times 4 \times 4$ maps and increasing to 70% for $64 \times 64 \times 64$ maps, as shown in Figure ??.




Figure 2.11: Profile of Computational Load of RRT in 3D

As such, it is safe to say that the bottleneck function for RRT is `textttedgeCollisions()`. This conclusion is strengthened by the fact that `edgeCollisions()` was implemented in the fastest possible way (without relying on approximations or implementing multithreading), whereas `findNearestConfig` was implemented without any optimizations (a possibility might have been K -nearest nodes). Finally, this conclusion supports prior research that collision detection takes up the vast majority of CPU execution time. As such, this is the function that was targeted for hardware acceleration.

Chapter 3

Motion Planning in Hardware

3.1 Defining the Collision Detection Unit

It was demonstrated in Section 2.3.2 that the critical function of RRT was edge collision detection. As such, the thesis proposes designing a functional unit that takes advantage of pipelining and parallelization to speed up the detection of edge collisions. Section 3.1.2 outlines the technical specifications for the functional unit. Section 3.1.3 outlines the performance specifications.

3.1.1 Edge Collision Function

Edge Collision Function Description: Edge collision function and algorithm, perhaps both for normal and parallelized?

3.1.2 Technical Specifications

Put simply, the functional unit that implements the edge collision detection function in hardware should have the same rough technical specifications as when the function is defined in software (Section 3.1.1). That is, it should take an edge e and an OGM and return a boolean value: True, if the edge collides with an obstacle, otherwise False. Table 3.1 outlines the required technical specifications for the functional unit.

Element	Description/Justification
Constraints	
Dimension N	N defines the dimension of the cubic configuration space for which the functional unit should take an identically sized OGM
Inputs	
Edge e	An Edge e defined for a 3D configuration space by two points $\{p1, p2\}$, each defined by a set of 3D coordinates $\{x, y, z\}$.
Space S	In an abstract sense, the edge collision detection function takes Space S as an input. In a more practical sense, the functional unit will take an $N \times N \times N$ Occupancy Grid Map
Control Inputs	The functional unit must also have ports for control signals such as clock, reset, start, etc. These are required for adding the unit to a processor.
Outputs	
Return Value	1 bit return value: 1 if collision, 0 otherwise.
Control Outputs	Output ports for control signals such as idle, done, ready, etc. These are required for adding the unit to a processor.

Table 3.1: Technical Specifications for Edge Collision Detection Unit

Improve Technical Specifications: More detail on control units.

3.1.3 Performance Specifications

Performance Specifications Functional Unit

3.2 HoneyBee

The Honey Bee has long been renowned for its tireless work ethic. But people rarely give the Honey Bee credit for its remarkable navigation and collision avoidance strategies during flight. As such, it is quite appropriate that this functional unit, designed to work tirelessly, rapidly and efficiently to execute collision detection computations, is named HoneyBee.

More Iterations of HoneyBee Design: Note: Currently this report only shows the design/build/measurement of the first pass at designing the functional unit (Designated HoneyBee-A, or HB-A). Final report will detail further iterations.

3.2.1 Design

HoneyBee-A Design

The first design iteration, designated HoneyBee-A (HB-A), was designed to take advantage of the performance improvements associated with pipelining. Figure 3.1 demonstrates how pipelining in hardware improves latency. By default, instructions are executed in order, one at a time. Pipelining takes advantage of operations that are independent of each other to reduce the number of clock cycles required to complete a set of instructions.




Figure 3.1: Pipelining to Improve Latency

Make my own version of this figure

3.2.2 Build

Hardware Description Languages

Introduction to Hardware Description Languages

High Level Synthesis

High Level Synthesis (HLS) is an automated hardware design process that takes design files (written in high-level languages, such as C, C++ or SystemC) specifying the algorithmic function of a piece of hardware, interprets those files and creates digital hardware designs that execute this function. It effectively translates programming languages into hardware description languages. Key advantages of using HLS is speed and verification. It is much faster and easier to define functionality in C than it is in a HDL such as Verilog, and thus design iterations are faster. It is also much simpler to verify one's design, as the functional units can be put through test benches written in C. This project used Vivado HLS to build the HoneyBee Unit.

HoneyBee-A Synthesis

Figure 3.2 shows the interface summary of successful synthesis of HoneyBee-A. Notice that the edge input has been split into 6 32 bit input ports.

RTL Ports	Dir	Bits	Protocol	Source Object	C Type
ap_clk	in	1	ap_ctrl_hs	honeybee	return value
ap_rst	in	1	ap_ctrl_hs	honeybee	return value
ap_start	in	1	ap_ctrl_hs	honeybee	return value
ap_done	out	1	ap_ctrl_hs	honeybee	return value
ap_idle	out	1	ap_ctrl_hs	honeybee	return value
ap_ready	out	1	ap_ctrl_hs	honeybee	return value
ap_return	out	32	ap_ctrl_hs	honeybee	return value
edge_p1_x	in	32	ap_none	edge_p1_x	scalar
edge_p1_y	in	32	ap_none	edge_p1_y	scalar
edge_p1_z	in	32	ap_none	edge_p1_z	scalar
edge_p2_x	in	32	ap_none	edge_p2_x	scalar
edge_p2_y	in	32	ap_none	edge_p2_y	scalar
edge_p2_z	in	32	ap_none	edge_p2_z	scalar

Figure 3.2: Interface Summary of HoneyBee-A Synthesis in Vivado HLS

Turn this into a table or get image from vivado

3.2.3 Measurement and Analysis

HoneyBee-A

The synthesis results of HoneyBee-A are shown in Table 3.2. It compares the execution time in microseconds for one edge to undergo collision detection if software and then in different synthesis “solutions”. MacOS and Ubuntu executing the function defined in honeybee.c have fairly similar results. Solution 1, which is the synthesized version of honeybee.c without any pipelining, was significantly slower. This is to be expected, as both MacOS and Ubuntu, operating on intel processors, would likely have some degree of pipelining and optimization of executing the compiled C code. However, significant improvements are observed once pipelining is implemented. Solutions 2-4 are increasing amounts of pipelining. Across the board, solutions 3 and 4 are roughly equal, but significantly faster than both solution 1 and the MacOS/Ubuntu execution times. Solution 4 shows a speedup of over 10x MacOS and Ubuntu.

Dimensions	Mac OS	Ubuntu	1	2	3	4
4x4x4	2	2	21.6	1.5	0.44	0.47
8x8x8	23	19	151	5.53	2.2	1.79
16x16x16	166	180	1133	41.37	13.08	12.11
32x32x32	1317	1424	8783	328	103	104

Table 3.2: Simulated performance of HB-A in microseconds

When HoneyBee-A is simulated in full RRT execution, we see similarly promising results. Table 3.3 shows the results of simulated RRT execution with HoneyBee-A. This is also shown in Figure 3.3.

Executions	K	Software	Sol. 1	Sol. 2	Sol. 3	Sol. 4
1221	100	0.420	100.724	0.400	0.125	0.126
2986	250	1.251	26.226	0.979	0.307	0.310
5719	500	1.997	50.229	1.875	0.589	0.594
8299	750	2.907	72.890	2.722	0.854	0.863
11148	1000	4.798	97.912	3.656	1.148	1.159
27203	2500	9.172	238.923	8.922	2.801	2.829
54499	5000	18.509	478.664	17.875	5.613	5.667
80952	7500	27.833	711.001	26.552	8.338	8.419
107487	10000	36.311	944.058	35.255	11.071	11.178

Table 3.3: RRT Simulated Execution Times with HB-A (seconds)




Figure 3.3: RRT Simulated Execution Time with HB-A (microseconds)

Make version of this chart in matplotlib for consistency and update axis

Expand Discussion of HoneyBee-A Results

Chapter 4

RISC-V Processor

4.1 Introduction to the Reduced Instruction Set Computer

4.1.1 Instruction Set Architecture

An Instruction Set Architecture can be thought of as an abstract model of a computer. On a broad level, it defines the data types, memory model, and registers of a computer, along with the instructions that it can execute.

It can also be thought as a “contract” between hardware and software developers. It is the promise made that the hardware will be able to execute all instructions defined in the ISA, and the limitation that software must be compiled into that set of instructions.

4.1.2 RISC Processor Design




Figure 4.1: 5-Stage RISC Datapath

4.2 RISC-V ISA

4.2.1 RV32I

The following is an excerpt from the RISC-V Specification, outlining the RV32I base integer instruction set [7]

RV32I was designed to be sufficient to form a compiler target and to support modern operating system environments. The ISA was also designed to reduce the hardware required in a minimal implementation. RV32I contains 40 unique instructions, though a simple implementation might ...[reduce] base instruction count to 38 total. RV32I can emulate almost any other ISA extension ...

Subsets of the base integer ISA might be useful for pedagogical purposes, but the base has been defined such that there should be little incentive to subset a real hardware implementation ...

Registers

RV32I defines 32 unprivileged registers, each 32 bits wide. They are designated **x0-x31**, where **x0** is a hard-wired value of 0, and registers **x1-x31** hold values that various instructions use. RISC-V uses the load-store method, meaning that all operations perform on two registers or a register and an immediate, rather than performing operations directly on memory addresses. In addition, the 33rd unprivileged register is the program counter **pc**. Table 4.1 shows the register state for the RV32I Base Integer Instruction Set.

Register	ABI Name	Description
x0	zero	Hard-wired zero
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5-7	t0-2	Temporaries
x8	s0/fp	Saved register/Frame pointer
x9	s1	Saved register
x10-11	a0-1	Function arguments/return values
x12-17	a2-7	Function arguments
x18-27	s2-11	Saved registers
x28-31	t3-6	Temporaries
pc	pc	Program counter

Table 4.1: Register State for RV32I Base Instruction Set

Instruction Formats

Table 4.2 demonstrates the format of each different instruction type.

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0			
funct7				rs2			rs1		funct3		rd			opcode		R-type	
imm[11:0]						rs1		funct3		rd			opcode		I-type		
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type	
imm[12]		imm[10:5]			rs2			rs1		funct3		imm[4:1]		imm[11]		opcode	B-type
imm[31:12]										rd				opcode		U-type	
imm[20]		imm[10:1]			imm[11]		imm[19:12]				rd			opcode		J-type	

Table 4.2: RV32I Base Instruction Formats

4.2.2 Motion Planning Extension

Motion Planning Extension: Full description of design of Non standard extension for motion planning. Should follow define, design, build, measure, analyse etc format.

4.3 PhilosophyV

Philosophy IV, written in 1903 by Mr. Owen Wister of the Class of 1882, recounts the antics of two Harvard students and their last minute attempts to study (or avoid studying) for an exam for which they are hopelessly unprepared. Similarly, this section details the process of my attempt to build a RISC-V processor, by far the most complex part of this Thesis, and a task for which I am unsure of my preparedness. As such, this processor is named Philosophy V; both in reference to the RISC-V ISA for which it is designed, and to the fact that my current situation seems much like a sequel to Mr. Wister's novel.

4.3.1 Baseline Implementation

Description of Baseline Philosophy V core

Figure 4.2 provides a schematic of the PhilosophyV processor.

Display bigger version of processor.

4.3.2 Implementing HoneyBee

Process of implementing honeybee into PhilV.




Figure 4.2: Philosophy V Processor

4.4 Performance Analysis

Comparative Performance Analysis of baseline and extended PhilV Core

Chapter 5

Conclusion

5.1 Discussion of Results

Discussion of Results

5.2 Evaluation of Success

Evaluation of Success

5.3 Future Work

Future Work

Bibliography

- [1] H. (Alexandrinus), *De gli automati, overo machine se moventi, Volume 2*.
- [2] V. I. B. U.-I. Isa, A. Waterman, Y. Lee, D. Patterson, K. Asanovi, and B. U.-I. Isa, “The RISC-V Instruction Set Manual v2.1,” *2012 IEEE International Conference on Industrial Technology, ICIT 2012, Proceedings*, vol. I, pp. 1–32, 2012.
- [3] N. Atay and B. Bayazit, “A motion planning processor on reconfigurable hardware,” in *Proceedings - IEEE International Conference on Robotics and Automation*, vol. 2006, pp. 125–132, 2006.
- [4] S. Murray, W. Floyd-Jones, G. Konidaris, and D. J. Sorin, “A Programmable Architecture for Robot Motion Planning Acceleration,” tech. rep.
- [5] G. S. Malik, K. Gupta, K. M. Krishna, and S. R. Chowdhury, “FPGA based combinatorial architecture for parallelizing RRT,” in *2015 European Conference on Mobile Robots, ECMR 2015 - Proceedings*, Institute of Electrical and Electronics Engineers Inc., nov 2015.
- [6] S. Murray, W. Floyd-Jones, Y. Qi, D. Sorin, G. Konidaris, and D. Robotics, “Robot Motion Planning on a Chip,” tech. rep.
- [7] A. Waterman, K. Asanovic, and SiFive Inc, “The RISC-V Instruction Set Manual,” vol. Volume I:, 2019.
- [8] C. E. Shannon, “XXII. Programming a computer for playing chess,” *The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science*, vol. 41, pp. 256–275, mar 1950.
- [9] M. Campbell, A. J. Hoane, and F. H. Hsu, “Deep Blue,” *Artificial Intelligence*, vol. 134, pp. 57–83, jan 2002.
- [10] S. M. LaValle, “Rapidly-Exploring Random Trees: A New Tool for Path Planning,” *In*, vol. 129, pp. 98–11, 1998.

- [11] S. M. LaValle and J. J. Kuffner, “Randomized kinodynamic planning,” *International Journal of Robotics Research*, vol. 20, pp. 378–400, may 2001.
- [12] Intel, “VTune Profiler,” 2019.
- [13] Thingbits, “LynxMotionHQuad500 Drone.”
- [14] RoboJackets, “RRT,” 2019.
<https://github.com/RoboJackets/rrt>.
- [15] M. Planning, “rrt-algorithms,” 2019.
<https://github.com/motion-planning/rrt-algorithms>.
- [16] Sourishg, “rrt-simulator,” 2017.
<https://github.com/sourishg/rrt-simulator>.
- [17] Vss2sn, “Path Planning,” 2019.
https://github.com/vss2sn/path_planning.
- [18] Olzhas, “RRT Toolbox,” 2017.

Glossary

RISC-V: (Pronounced “risk-five”) is an open-source and extendible ISA developed by the University of California, Berkeley. It is established on the principles of a RISC, a class of instruction sets that allow a processor to have fewer CPI than a CISC.

A priori: relating to or denoting reasoning or knowledge which proceeds from theoretical deduction rather than from observation or experience..

Automata: moving mechanical devices made in imitation of human beings.

Configuration: A specification of a robots location, position, and setting in a space. For example, when a robot is represented by a single point in 3D space, its configuration is merely x, y, and z coordinates. But if a robot is represented as a 3D humanoid with a head, body, arms and legs, then its configuration would be its position in 3D space, its orientation in 3D, and the position of all its joints and limbs such that the space being taken up by the robot can be exactly determined. It is obvious then that as the physical complexity of a robot increases, so too does the complexity of representing it algorithmically..

Dijkstra’s algorithm: An algorithm for finding the shortest path between two nodes in a graph..

Hardware acceleration: The process of speeding up the execution of a function by implementing part or whole of that function specifically in hardware..

Mathematically complete: An algorithm is mathematically complete if it will always find all solutions..

Probabilistically complete: Describes an algorithm with a likelihood of finding a solution that approaches one as its runtime approaches infinity..

Real-time: Describes a system in which input data is processed in such a time period such that it is available almost immediately. Systems such as missile guidance or collision avoidance in cars are an example of real-time systems.

Workspace: The space which a robot and obstacles occupy in motion planning problems.

Advanced RISC Machine (ARM): A family of Reduced Instruction Set Computing architectures for computer processors, configured for various environments..

Application Specific Processor (ASP): A computer processor that has been optimized for a specific function or set of functions that support a given application..

Complex Instruction Set Computer (CISC): A computer in which single instructions can execute several low-level operations (such as a load from memory, an arithmetic operation, and a memory store) or are capable of multi-step operations or addressing modes within single instructions..

Degree-of-Freedom (DOF): Refers to the number of independent factors that describe the configurations in which a robot can exist and motion can occur..

Field Programmable Gate Array (FPGA): An integrated circuit designed to be configured by a designer after manufacturing – hence the term ‘field-programmable’. Its behaviour is specified in software, using a HDL.

Hardware Description Language (HDL): A computer language used for designing computer hardware. It is used to define the behaviour of modules, simulate their performance, and synthesize them on an FPGA.

High Level Synthesis (HLS): An automated hardware design process that takes software written in high-level languages (often C, C++) that algorithmically defines a function, and converts that into HDL that implements that function..

Instruction Set Architecture (ISA): An abstract model of a computer, that defines the instructions, registers, memory behaviour, and other attributes of a computer architecture. It can be thought of as the contract between software and hardware developers, as the ISA lists the instructions that software may be implemented in, and the instructions that a processor must support..

Occupancy Grid Map (OGM): A method of representing a 2D or 3D space by dividing it into discrete grids and marking each whole grid as ‘occupied’, even if only part of it is..

Probabilistic Road Map (PRM): A motion planning algorithm that randomly samples free space, and then connects sampled configurations with nearby configurations to build a map..

Reduced Instruction Set Computer (RISC): A computer architecture based on a small number of instructions executed in a small number of cycles..

Rapidly-exploring Random Tree (RRT): an algorithm designed to efficiently search, and thus plan a path through, a high-complexity environment by randomly sampling points and building a tree. The algorithm randomly samples points, draws an edge from the nearest currently existing node in the tree, to grow the tree in the space..

Real-Time Operating Systems (RTOS): A type of operating system designed to operate on inputs as they come in, without buffer delays..

RISC-V 32-Bit Integer (RV32I): One of the four base ISAs within RISC-V. While it implements integer values only in 32-bit representations, it contains the minimal number of instructions for a fully working computer processor..

System on Chip (SoC): A chip that includes all required components of a working computer, including one or more CPUs, memory, I/O ports, etc..

Unmanned Aerial Vehicle (UAV): An aircraft without a human pilot on board. It may be piloted remotely completely by a human pilot, autonomously pilot itself, or a mixture of the two..

Appendices

Appendix A

Project Repository

This project's repository can be found at github.com/AnthonyKenny98/Thesis and contains multiple subrepositories. It has the following structure.

Research

This folder holds the academic papers that constitute the background research of this Thesis.

Writeups

This folder holds the writeups required for this Thesis, including checkpoints in fulfillment of Harvard's ES100hf class and this Final Report

RRT

github.com/AnthonyKenny98/RRT

This subrepository holds both the 2D and 3D implementations of RRT used for this thesis, along with the tools required for both VTune Profiler and internal timing analysis.

HoneyBee

github.com/AnthonyKenny98/HoneyBee

This subrepository holds the HoneyBee functional unit, a hardware implementation of collision detection.

PhilosophyV

github.com/AnthonyKenny98/PhilosophyV

This subrepository holds the PhilosophyV RISC-V chip

Appendix B

Budget

Budget

Appendix C

RRT Supporting Documentation

C.1 Justification of Modelling UAV as Prism

While it is possible for a UAV to be modelled in precise detail, taking into account its exact shape, more often UAVs are modelled as a 3D prism in motion planning problems, for the following reasons:

- It is a rare case that the negative space gained by modelling in such detail is utilised
- Representation of drone's configuration is much more complex.
- Computing edge collisions is much more computationally intensive.




Figure C.1: Modelling a UAV as a Rectangular Prism. Red highlight demonstrates the model, overlayed over the exact schematic. Figure C.1a shows how a drone can be modelled in high detail, but gains little useful free space when compared with Figure C.1b, which models a drone as a rectangular prism.[13]

C.2 Full Technical Specifications for RRT Implementation

General Specifications	
Requirement	Description and Justification
Implemented in C/C++	As outlined in Section 1.3.3, the critical step in determining the design of specialized hardware to accelerate RRT is CPU performance analysis of the algorithm to determine computational hot-spots. Implementations in C allow for the use of certain CPU profiling tools, unlike higher-level languages such as Python.
3D Workspace	The computational requirements of RRT in 3D differs somewhat to that in 2D. Since autonomous UAVs operate in 3D space, it was necessary to have a 3D implementation to analyse.
UAV modelled in 3D as a rectangular prism	In theory, it is possible to model a UAV much more precisely than a rectangular prism, taking into account its shape and negative space. However, in reality, modelling a UAV as a 3D rectangular prism, defined by coordinates $\{x, y, z\}$ and Euler angles $\{\alpha, \beta, \gamma\}$, is more than sufficient (and more efficient). See Appendix C.1 for justification of this.
Mathematically Complete Collision Detection	When RRT is implemented for educational purposes, the edge collision calculations are often simplified to a sampling model which is probabilistically complete but not mathematically complete. In other words, it will catch most collisions by sampling a number of points along each edge, but there is always a possibility of an undetected collision. In real world applications, collisions must be calculated by method of geometric intersection to ensure all collisions are detected.

Table C.1: General Technical Specifications for RRT Implementation

Required Parameters	
Parameter	Description and Justification
ϵ (a.k.a. Δq)	The maximum difference between two configurations. Larger values of ϵ can solve less obstacle dense problems faster, but take longer to solve problems with tight corners.
K	The maximum number of configurations. This is largely correlative to the amount of time the user will allow the algorithm to run. Larger values of K will take longer but generate better paths, while smaller values will execute for less time but generate more jagged paths or may not reach the goal node. The value of K was varied to find the minimum execution time while still reaching the goal with high probability.
DIM	The upper bound of each axis of a $DIM \times DIM \times DIM$ Workspace. Larger values leave more space to be explored, and thus require larger values of K to reach the goal with high likelihood.
Goal Bias	The given probability that the graph will extend the graph ϵ distance from an existing configuration to a new configuration in the direction of the goal.

Table C.2: Required Parameters for RRT Implementation

C.3 Assessment of Existing RRT Implementations

Repository	General Requirements				Parameters			
	Language	2D/3D	Object Model	Collision Detect	ϵ	K	DIM	Goal Bias
RoboJackets[14]	C++	2D	Point	Complete	Yes	Yes	Yes	No
Motion-Planning[15]	Python	ND	Point	Incomplete	Yes	Yes	No	Yes
Sourishg[16]	C++	2D	Point	Incomplete	Yes	Yes	No	No
Vss2sn[17]	C++	2D	Point	Complete	Yes	Yes	No	No
olzhas[18]	Matlab	2D	Point	Complete	Yes	Yes	No	No

Table C.3: Evaluation of Existing Open-Source Implementations of RRT. Links to Github repositories can be found in the Bibliography.

C.4 Implementation of Key RRT Functions

Algorithm C.1: `getRandomConfig()` as implemented for RRT

Inputs: Dimensionality N ,
 Upper Axis Bound DIM
Output: Random Configuration q
 $q.x \leftarrow \text{randomFloat}(DIM)$
 $q.y \leftarrow \text{randomFloat}(DIM)$
 $q.\alpha \leftarrow \text{randomFloat}(2\pi)$
if $N == 3$ **then**
 $q.z \leftarrow \text{randomFloat}(DIM)$
 $q.\beta \leftarrow \text{randomFloat}(2\pi)$
 $q.\gamma \leftarrow \text{randomFloat}(2\pi)$
end
return q ;

Where `randomFloat(max)` returns a float between 0 and `max`.

Algorithm C.2: `findNearestConfig()` as implemented for RRT

Inputs: Graph G ,
 New Configuration q_{new}
Output: Nearest Configuration $q_{nearest}$
 $q_{nearest} \leftarrow G.q_{init}$
for $k = 0$ to $G.\text{existing_nodes}$ **do**
 if $\text{distance}(q_{new}, G.q[k]) < \text{distance}(q_{new}, q_{nearest})$ **then**
 $q_{nearest} \leftarrow G.q[k]$
 end
end
return $q_{nearest}$

Where `distance(q_1 , q_2)` returns the Euclidean distance between two configurations.

Algorithm C.3: stepFromNearest() as implemented for RRT

Inputs: Configuration in Graph $q_{nearest}$,
 New Configuration q_{new} ,
 Goal Bias B ,
 Maximum Step Distance ϵ ,
 Graph G

Output: Updated New Configuration q_{new}

```

if distance( $q_{nearest}$ ,  $q_{new}$ ) >  $\epsilon$  then
  | if randomFloat(1) <  $B$  then
  | |  $q_{new} \leftarrow \text{stepTowardConfig}(q_{nearest}, G.q_{goal})$ 
  | end
  | else
  | |  $q_{new} \leftarrow \text{stepTowardConfig}(q_{nearest}, q_{new})$ 
  | end
end
return  $q_{new}$ ;
  
```

Where $\text{stepTowardConfig}(q_1, q_2)$ returns a configuration ϵ from q_1 in the direction of q_2 .

Algorithm C.4: configCollision() as implemented for RRT

Inputs: Dimensionality N ,
 Occupancy Grid Map (N -Dimensional Array) O ,
 Configuration q

Output: Boolean

```

if  $N == 2$  then
  | return  $O[\text{gridLookup}(q.x)][\text{gridLookup}(q.y)]$ 
end
else
  | return  $O[\text{gridLookup}(q.x)][\text{gridLookup}(q.y)][\text{gridLookup}(q.z)]$ 
end
  
```

Where O is a N -Dimensional array of booleans, with True representing an occupied grid and false representing an unoccupied one. $\text{gridLookup}()$ is a function that maps a floating point coordinate to the correct integer of the grid in which it resides. For a map resolution of one, this is as simple as rounding a float down to an integer.

While seemingly complex, the above algorithm merely steps through the mathematical process of checking the relevant x , y , and z planes for a point of intersection with the

Algorithm C.5: `configCollision()` as implemented for RRT for 3D

```

Inputs:   Edge  $e$ ,
             Occupancy Grid Map (3-Dimensional Array)  $O$ ,
             Maximum Step Distance  $\epsilon$ 

Output:  Boolean

 $q_{min} \leftarrow \text{minConfig}(e.q_1, e.q_2)$ 
for ( $x = q_{min}.x$  to  $q_{min}.x + \epsilon$ ) do
     $q_{intersection} \leftarrow \text{edgeIntersectsPlane}(e, x)$ 
    if  $O[q_{intersection}.x][q_{intersection}.y][q_{intersection}.z]$  then
        return true
    end
end
for ( $y = q_{min}.y$  to  $q_{min}.y + \epsilon$ ) do
     $q_{intersection} \leftarrow \text{edgeIntersectsPlane}(e, y)$ 
    if  $O[q_{intersection}.x][q_{intersection}.y][q_{intersection}.z]$  then
        return true
    end
end
for ( $z = q_{min}.z$  to  $q_{min}.z + \epsilon$ ) do
     $q_{intersection} \leftarrow \text{edgeIntersectsPlane}(e, z)$ 
    if  $O[q_{intersection}.x][q_{intersection}.y][q_{intersection}.z]$  then
        return true
    end
end
return false

```

edge e . It then looks up the OGM O to see if the grid corresponding with the point of intersection is occupied. If so, then it reports a collision by returning false. The function `edgeIntersectsPlane` follows the geometrical process of detecting a segment-plane intersection outlined in Appendix C.5. q_{min} is calculated to be the origin point of the grid closest to the origin. In other words, the algorithm does not check for intersections throughout the entire map, only the maximum number of grids that could possibly be intersected by the edge e , given the location of the two points of the edge, $e.p_1$ and $e.p_2$, and the maximum edge length ϵ . The algorithm for `edgeCollision` in 2D can be inferred from the above, instead checking segment-line intersections for x and y lines.

C.5 Geometrically Determining Segment-Plane Intersection

The method of edge collision detection in this project's implementation of RRT relies on detecting segment-plane intersections. The planes are always set up to be parallel with either the $x - y$ plane, the $x - z$ plane, or the $y - z$ plane. Figure C.2 demonstrates this point.




Figure C.2: Using Parallel Planes to determine Edge Collisions with Grids

A plane can be defined by 3 points, P_a , P_b , and P_c . In practice, the points defining a plane parallel to the x - y plane would have the following points:

$$P_a = (x, y, z)$$

$$P_b = (x + \Delta x, y, z)$$

$$P_c = (x, y + \Delta y, z)$$

Two vectors, \vec{AB} and \vec{AC} can be determined.

The normal to the plane is the cross product:

$$\vec{AB} \times \vec{AC}$$

And the equation of the plane written as:

$$a(x - x_0) + b(y - y_0) + c(z - z_0) = 0$$

$$ax + by + cz = ax_0 + by_0 + cz_0$$

Where $\langle a, b, c \rangle$ is the normal to the plane and (x_0, y_0, z_0) is one of the points P_a , P_b , or P_c . The RHS can be set to equal d , leaving:

$$ax + by + cz = d$$

Now, the equation of a line can be written in the form:

$$ax + by + cz = 0$$

And can be parameterized in the following form:

$$\begin{cases} x = x_1 + t(x_2 - x_1) \\ y = y_1 + t(y_2 - y_1) \\ z = z_1 + t(z_2 - z_1) \end{cases}$$

To find the point of intersection, we substitute the equation of the line into the equation of the plane, yielding:

$$a(x_1 + t(x_2 - x_1)) + b(y_1 + t(y_2 - y_1)) + c(z_1 + t(z_2 - z_1)) = d$$

Rearranging to find an expression for t :

$$t = \frac{d - (ax_1 + by_1 + cz_1)}{a(x_2 - x_1) + b(y_2 - y_1) + c(z_2 - z_1)}$$

Knowing t , we can find the point of intersection, P_X to be:

$$\begin{cases} x_X(t) = x_1 + t(x_2 - x_1) \\ y_X(t) = y_1 + t(y_2 - y_1) \\ z_X(t) = z_1 + t(z_2 - z_1) \end{cases}$$

Finally, the following equalities are evaluated to see if the point lies on the segment:

$$x_1 \leq x_X \leq x_2$$

$$y_1 \leq y_X \leq y_2$$

$$z_1 \leq z_X \leq z_2$$

If so, then the grids corresponding to the point of intersection can be marked as intersected.

C.6 Timing Methodology of RRT Analysis

VTune Profiler

VTune Profiler performance profiler is an application for software performance analysis. It provides functionality to examine hot-spots for CPU execution time through a top down analysis. The top down analysis tool shows the percentage of CPU time taken up by each function. It was used to initially profile the algorithm's performance.

Internal Timing

There are several limitations to using VTune Profiler. First, it can only profile software running on Intel processors, which implement the x86 ISA. In anticipation of potentially needing to run performance analysis on a RISC-V processor, another method would have been required. Secondly, VTune Profiler takes a long time to run, as it needs to conduct much analysis that is extraneous to the purpose of this thesis. This became prohibitive when it came to conducting hundreds of tests for different parameterization, with each test running RRT a minimum of 100 times. Finally, it was not customizable to ignore certain parts of the implementation which were necessary for logging. While the implementation was designed in such a way that these should not interfere, it led to a lot of irrelevant data. A simple and effective alternative for measuring execution performance was to insert timing functionality into the software itself.

Internal timing was implemented based on the inbuilt C `clock()` function and `CLOCKS_PER_CYCLE` macro, and wrapping each function of interest in a performance tracking struct. This can be seen in the project's RRT sub-repository under `performance.h`.

Comparison

Before proceeding to use the internal timing method to profile the software implementation of RRT, it was important to verify that the two methods yielded similar results for the same program. Table C.4 summarizes the results of analysis of a simple C executable. The program calls 5 functions, $\{A, B, C, D, E\}$, each a simple iteration in which an integer is incremented. Since the Internal Timing method returned similar results to the (trusted) VTune Profiler, it was considered to be a reliable method. While it was encouraging to see both methods returned similar results for absolute execution time, the more important metric was the similarity in percentage of total execution time. For good measure, a χ^2 test of hypothesis was conducted and for one degree of freedom showed more than acceptable results.

function	Vtune Profiler		Internal Timing		χ^2
	time (s)	time (% total)	time (s)	time (% total)	
A	0.488	57.4%	0.497	57.6%	0.00016
B	0.2	23.5%	0.198	23.1%	0.00002
C	0.102	12.0%	0.099	11.5%	0.00009
D	0.048	5.7%	0.049	5.6%	0.00002
E	0.012	1.4%	0.019	2.2%	0.00408

Table C.4: Comparison of Timing Methods

C.7 Justification of K:DIM Ratio

Make all reference to "Drone" uniform to "UAV" (this will require changing the "a" to "an" before each occurrence - wait actually, do I?.) Also maybe I should clarify this to rotary wing UAVs.

Spell check entire doc

Make all figure headings bold with subtitles.

Todo list

<input type="checkbox"/>	Fix Cover Page Formatting	i
<input type="checkbox"/>	Rewrite Abstract	i
<input type="checkbox"/>	Define goal	1
<input type="checkbox"/>	cite	1
<input type="checkbox"/>	cite	2
<input type="checkbox"/>	Discuss Moore's law and denard scaling	3
<input type="checkbox"/>	Improve Problem Statement	4
<input type="checkbox"/>	Revise End User	4
<input type="checkbox"/>	Project Specifications	9
<input type="checkbox"/>	Summary of Results	10
<input type="checkbox"/>	Update once I have properly defined goals and objectives	11
<input type="checkbox"/>	This is all really ugly and should be explained better	15
<input type="checkbox"/>	Redo this diagram	16
<input type="checkbox"/>	System Diagram Here	19
<input type="checkbox"/>	Edge Collision Function Description	27
<input type="checkbox"/>	Improve Technical Specifications	28
<input type="checkbox"/>	Performance Specifications Functional Unit	28
<input type="checkbox"/>	More Iterations of HoneyBee Design	29
<input type="checkbox"/>	Make my own version of this figure	29
<input type="checkbox"/>	Introduction to Hardware Description Languages	30
<input type="checkbox"/>	Turn this into a table or get image from vivado	31
<input type="checkbox"/>	Make version of this chart in matplotlib for consistency and update axis	33
<input type="checkbox"/>	Expand Discussion of HoneyBee-A Results	33
<input type="checkbox"/>	Motion Planning Extension	36
<input type="checkbox"/>	Description of Baseline Philosophy V core	36
<input type="checkbox"/>	Display bigger version of processor.	36
<input type="checkbox"/>	Process of implementing honeybee into PhilV.	36
<input type="checkbox"/>	Comparative Performance Analysis of baseline and extended PhilV Core	38
<input type="checkbox"/>	Discussion of Results	39
<input type="checkbox"/>	Evaluation of Success	39
<input type="checkbox"/>	Future Work	39

<input type="checkbox"/>	Budget	47
<input type="checkbox"/>	Make all reference to "Drone" uniform to "UAV" (this will require changing the "a" to "an" before each occurrence - wait actually, do I?.) Also maybe I should clarify this to rotary wing UAVs.	58
<input type="checkbox"/>	Spell check entire doc	58
<input type="checkbox"/>	Make all figure headings bold with subtitles.	58